

```
import numpy as np
import pandas as pd
from scipy import stats
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Read Dataset
```

```
df = pd.read_csv('./income_evaluation.csv')
df.columns = list(map(lambda a: a.lstrip(), df.columns))
df.head()
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	sex	marital-status	occupation	relationship	race
0	Male	Never-married	Adm-clerical	Not-in-family	White
1	Male	Married-civ-spouse	Exec-managerial	Husband	White
2	Male	Divorced	Handlers-cleaners	Not-in-family	White
3	Male	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	Female	Married-civ-spouse	Prof-specialty	Wife	Black

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

Exploratory Data Analysis

Count Missing Values

```
missing_count = {}
for column in df.columns:
    df[column].replace(' ?', np.NaN, inplace=True)
    missing_count[column] = df[column].isna().sum()

for key in missing_count:
    print(f"{key}: {missing_count[key]}")
```

```
age: 0
workclass: 1836
fnlwgt: 0
education: 0
education-num: 0
marital-status: 0
occupation: 1843
relationship: 0
race: 0
sex: 0
capital-gain: 0
capital-loss: 0
hours-per-week: 0
native-country: 583
income: 0
```

```
df.isnull().sum()
```

```
age          0
workclass    1836
fnlwgt       0
education    0
education-num 0
marital-status 0
occupation   1843
relationship 0
race         0
sex          0
capital-gain 0
capital-loss 0
hours-per-week 0
native-country 583
income       0
dtype: int64
```

Count Numerical & Categorical Features

```
numerical_features=[features for features in df.columns if
df[features].dtype!='O']
categorical_features = [features for features in df.columns if
```

```
df[features].dtype == 'O']

print("Numerical Features: ", ', '.join(numerical_features), "\n\
Categorical Features: ", ', '.join(categorical_features))
```

Numerical Features: age, fnlwgt, education-num, capital-gain,
capital-loss, hours-per-week
Categorical Features: workclass, education, marital-status,
occupation, relationship, race, sex, native-country, income

Continuous and Discrete Features

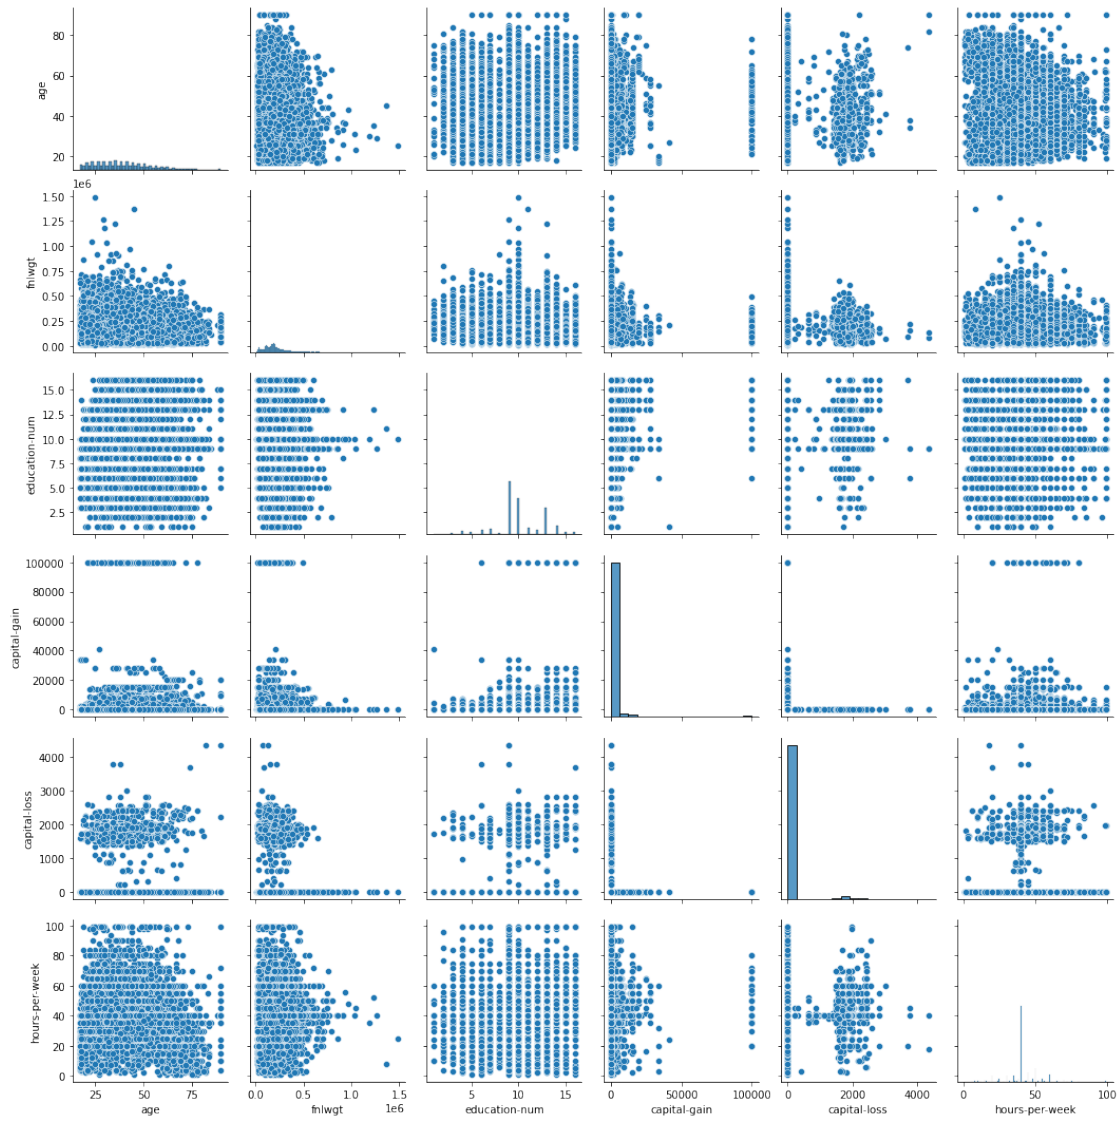
```
continuous_features=[features for features in numerical_features if  
len(pd.unique(df[features]))>25]  
discrete_features=[features for features in numerical_features if  
len(pd.unique(df[features]))<=25]  
print("Continuous Features: ", ', '.join(continuous_features), "\n\
Discrete Features: ", ', '.join(discrete_features))
```

Continuous Features: age, fnlwgt, capital-gain, capital-loss, hours-
per-week
Discrete Features: education-num

Pair Plots



```
sns.pairplot(df,kind='scatter')



<seaborn.axisgrid.PairGrid at 0x7f2e790cdf90>
```





Histograms

```
for col in numerical_features:
    sns.histplot(df[col])
plt.show()
```















Displot



```
for col in df.columns:  
    sns.displot(df[col], kde=True, legend=True,  
aspect=2).set_xticklabels(rotation=90)
```


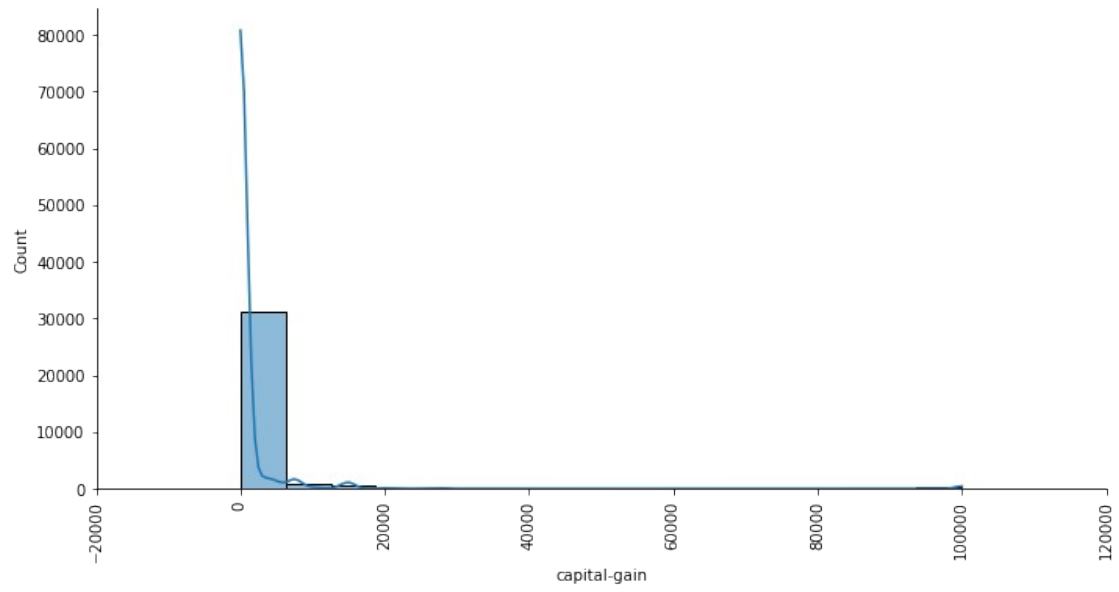

























Are there any outliers present?

```
for col in numerical_features:  
    df.boxplot(col)  
    plt.show()
```









Feature Engineering

Replace none values

```
df_modified=df
for col in numerical_features:
    df_modified[col].replace( np.NaN, df.median(axis=0,skipna=True)
[col],inplace=True)
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3:
FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

This is separate from the ipykernel package so we can avoid doing imports until

```
for col in categorical_features:
    df_modified[col].replace( np.NaN, df.mode()[col][0],inplace=True)
```

Replace categorical features with numerical values

```
df_cat_to_num=pd.get_dummies(df_modified)
```

```
df_cat_to_num.head()
```

	age	fnlwgt	education-num	capital-gain	capital-loss	\
0	5.130679	266.098568	13	0.559558	0.274534	

1	5.612071	274.345493	13	0.559558	0.274534
2	5.081547	410.077483	9	0.559558	0.274534
3	5.728022	425.012149	7	0.559558	0.274534
4	4.520327	495.934871	13	0.559558	0.274534

	hours-per-week	workclass_ Federal-gov	workclass_ Local-gov	\
0	39.163404	0	0	
1	31.621057	0	0	
2	39.163404	0	0	
3	39.163404	0	0	
4	39.163404	0	0	

	workclass_ Never-worked	workclass_ Private	...	native-country_ Scotland	\
0	0	0	...		
0					
1	0	0	...		
0					
2	0	1	...		
0					
3	0	1	...		
0					
4	0	1	...		
0					

	native-country_ South	native-country_ Taiwan	native-country_ Thailand	\
0	0	0		
0				
1	0	0		
0				
2	0	0		
0				
3	0	0		
0				
4	0	0		
0				

	native-country_ Trinidad&Tobago	native-country_ United-States	\
0	0	1	
1	0	1	
2	0	1	
3	0	1	
4	0	0	

	native-country_ Vietnam	native-country_ Yugoslavia	income_ <=50K	\
0	0	0	1	
1	0	0	1	


2	0	0	1
3	0	0	1
4	0	0	1




	income_ >50K
0	0
1	0
2	0
3	0
4	0


[5 rows x 107 columns]

Using boxcox to normalise data

```
df_normalized=df_cat_to_num
df_normalized['capital-gain']=df_normalized['capital-gain']+5
df_normalized['capital-loss']=df_normalized['capital-loss']+5
for feature in continuous_features:
    df_normalized[feature] =
pd.DataFrame(stats.boxcox(df_normalized[feature])[0])
# print(df[feature])
sns.displot(df_normalized[feature], kde=True, legend=True,
height=5, aspect=2)
```







Removing outliers

```
for feature in numerical_features:
    q75,q25 = np.percentile(df_normalized.loc[:,feature],[75,25])
    intr_qr = q75-q25

    max_range= q75+(1.5*intr_qr)
    min_range = q25-(1.5*intr_qr)

    df_normalized[feature].values[df_normalized[feature].values <
min_range] = min_range
    df_normalized[feature].values[df_normalized[feature].values >
max_range] = max_range

df_normalized.head()
```

	age	fnlwgt	education-num	capital-gain	capital-loss	\
0	5.130679	266.098568	13	0.559558	0.274534	
1	5.612071	274.345493	13	0.559558	0.274534	
2	5.081547	410.077483	9	0.559558	0.274534	
3	5.728022	425.012149	7	0.559558	0.274534	
4	4.520327	495.934871	13	0.559558	0.274534	

	hours-per-week	workclass_ Federal-gov	workclass_ Local-gov	\
0	39.163404	0	0	
1	31.621057	0	0	
2	39.163404	0	0	
3	39.163404	0	0	
4	39.163404	0	0	

	workclass_ Never-worked	workclass_ Private	...	native-country_ Scotland	\
0		0	0	...	
0					

1	0	0	...
0			
2	0	1	...
0			
3	0	1	...
0			
4	0	1	...
0			

	native-country_ South Thailand \	native-country_ Taiwan	native-country_
0	0	0	
0			
1	0	0	
0			
2	0	0	
0			
3	0	0	
0			
4	0	0	
0			



	native-country_ Trinidad&Tobago	native-country_ United-States \
0	0	1
1	0	1
2	0	1
3	0	1
4	0	0



	native-country_ Vietnam	native-country_ Yugoslavia	income_ <=50K
\			
0	0	0	1
1	0	0	1
2	0	0	1
3	0	0	1
4	0	0	1


	income_ >50K
0	0
1	0
2	0
3	0
4	0

```
[5 rows x 107 columns]
```

```
for col in numerical_features:  
    df_normalized.boxplot(col)  
    plt.show()
```







Correlation

```
corr=df_normalized.corr()['income_ <=50K'][:]  
corr
```

age	-0.264097
fnlwgt	0.004985
education-num	-0.339579
capital-gain	NaN
capital-loss	NaN

```

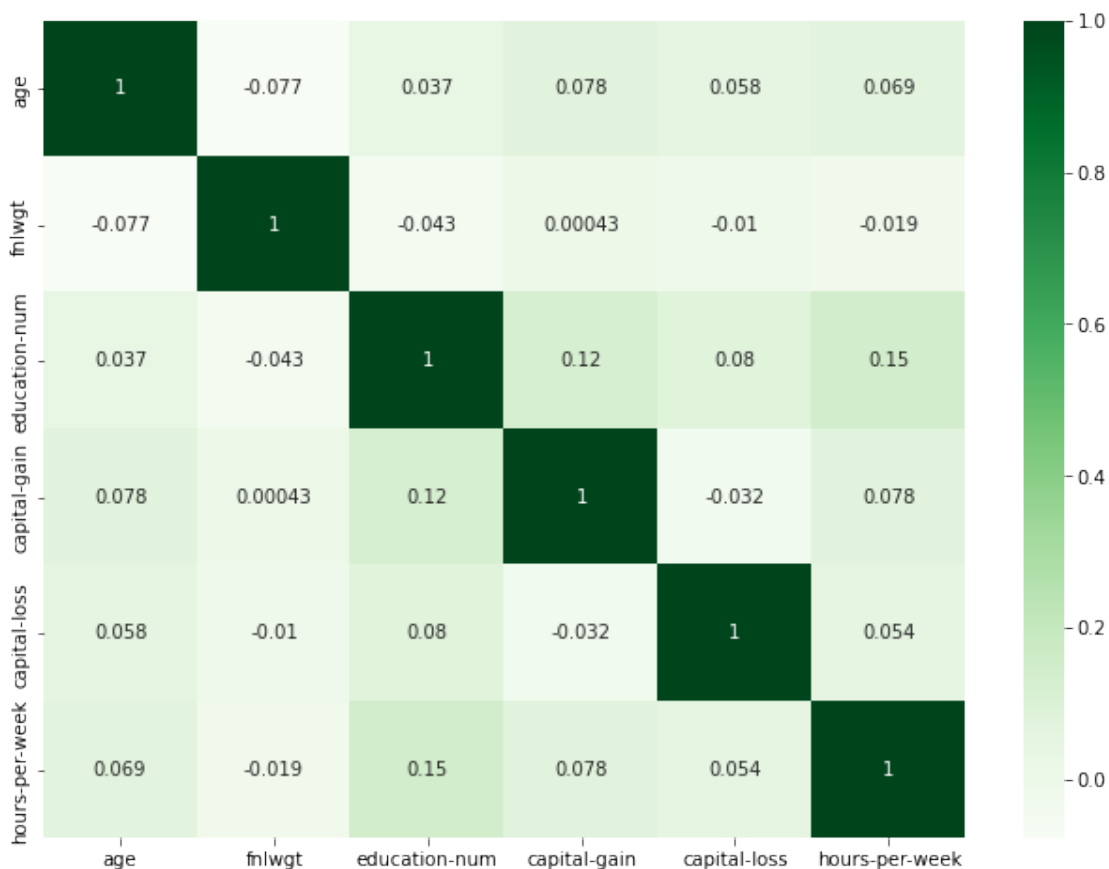
native-country_ United-States -0.038942
native-country_ Vietnam      0.017649
native-country_ Yugoslavia   -0.006959
income_ <=50K                1.000000
income_ >50K                 -1.000000
Name: income_ <=50K, Length: 107, dtype: float64

```

```

plt.figure(figsize=(11,8))
sns.heatmap(df_modified.corr(), cmap="Greens",annot=True)
plt.show()

```



##Modelling

```

X = df_normalized.drop(columns=['income_ >50K', 'income_ <=50K'])
y = df_normalized['income_ >50K']

```

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3)

```

Logistic regression tuning

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

```

```

score_list = []
solvers={'newton-cg':['l2', 'none'], 'lbfgs': ['l2',
'none'], 'liblinear':['l1', 'l2'], 'sag':['l2', 'none'], 'saga':
['elasticnet', 'l1', 'l2', 'none']}
for solver, penalty in solvers.items():
    for p in penalty:
        try:
            logistic_regression = LogisticRegression(penalty=p,
solver=solver)

            logistic_regression.fit(X_train, y_train)
            y_predict = logistic_regression.predict(X_test)
            accuracy = accuracy_score(y_test, y_predict)
            score_list.append([p, solver, accuracy])
        except:
            print("an error has occurred")

df_LR = pd.DataFrame(score_list,
columns=['penalty', 'solver', 'accuracy'])
print(df_LR)

```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/optimize.py:212:
ConvergenceWarning: newton-cg failed to converge. Increase the number
of iterations.

ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.
py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as
shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.
py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as
shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_sag.py:35

```
4: ConvergenceWarning: The max_iter was reached which means the coef_
did not converge
  ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_sag.py:35
4: ConvergenceWarning: The max_iter was reached which means the coef_
did not converge
  ConvergenceWarning,

an error has occurred
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_
_sag.py:354: ConvergenceWarning: The max_iter was reached which means
the coef_ did not converge
  ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_sag.py:35
4: ConvergenceWarning: The max_iter was reached which means the coef_
did not converge
  ConvergenceWarning,
```

	penalty	solver	accuracy
0	l2	newton-cg	0.834579
1	none	newton-cg	0.833965
2	l2	lbfgs	0.827413
3	none	lbfgs	0.826697
4	l1	liblinear	0.834783
5	l2	liblinear	0.834579
6	l2	sag	0.814515
7	none	sag	0.814515
8	l1	saga	0.806428
9	l2	saga	0.807043
10	none	saga	0.806428

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_
_sag.py:354: ConvergenceWarning: The max_iter was reached which means
the coef_ did not converge
  ConvergenceWarning,
```

```
score_list_c=[]
for c in [100, 10, 1.0, 0.1, 0.01]:
    logistic_regression = LogisticRegression(solver='liblinear',
penalty='l1', C=c)
    logistic_regression.fit(X_train, y_train)
    y_predict = logistic_regression.predict(X_test)
    accuracy = accuracy_score(y_test, y_predict)
    score_list_c.append([c, accuracy])
df_LR_1 = pd.DataFrame(score_list_c, columns=['c', 'accuracy'])
print(df_LR_1)
```

	c	accuracy
0	100.00	0.834374
1	10.00	0.834272

```
2    1.00  0.835091
3    0.10  0.832736
4    0.01  0.820555
```

```
final_lr_model=LogisticRegression(solver='liblinear', penalty='l1',
C=1.0)
final_lr_model.fit(X_train, y_train)
y_predict = final_lr_model.predict(X_test)
accuracy = accuracy_score(y_test,y_predict)
accuracy

0.8350905926911659
```

Random forest classifier tuning

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import make_classification
score_list_rfc=[]
for depth in range(1,40):
    rfc= RandomForestClassifier(random_state=0,max_depth=depth)
    rfc.fit(X_train,y_train)
    y_predict = rfc.predict(X_test)
    accuracy = accuracy_score(y_test,y_predict)
    score_list_rfc.append([depth,accuracy])
```

```
score_list_rfc
```

```
[[1, 0.7648684614597195],
 [2, 0.7648684614597195],
 [3, 0.7968062237690654],
 [4, 0.8125703756781656],
 [5, 0.8257754120176067],
 [6, 0.8264919643771113],
 [7, 0.8291534445695568],
 [8, 0.8288463507011977],
 [9, 0.8302794554202068],
 [10, 0.8327362063670796],
 [11, 0.8338622172177296],
 [12, 0.8365236974101751],
 [13, 0.8386733544886887],
 [14, 0.8393899068481933],
 [15, 0.8399017299621251],
 [16, 0.8384686252431159],
 [17, 0.8421537516634251],
 [18, 0.8404135530760569],
 [19, 0.8376497082608251],
 [20, 0.8391851776026206],
 [21, 0.8376497082608251],
 [22, 0.8376497082608251],
 [23, 0.8383662606203296],
 [24, 0.8353976865595251],
 [25, 0.8353976865595251],
```

```
[26, 0.8355000511823114],
[27, 0.8328385709898659],
[28, 0.8340669464633023],
[29, 0.8334527587265841],
[30, 0.8314054662708568],
[31, 0.8309960077797113],
[32, 0.8299723615518477],
[33, 0.8282321629644794],
[34, 0.8289487153239841],
[35, 0.8265943289998976],
[36, 0.8280274337189067],
[37, 0.8276179752277613],
[38, 0.8300747261746341],
[39, 0.8284368922100522]]
```

accuracy

0.8421537516634251

```
score_list_rfc_1=[]
for trees in range(100,1001,100):
    rfc=
    RandomForestClassifier(random_state=0,max_depth=17,n_estimators=trees)
    rfc.fit(X_train,y_train)
    y_predict = rfc.predict(X_test)
    accuracy = accuracy_score(y_test,y_predict)
    score_list_rfc_1.append([trees,accuracy])
```

score_list_rfc_1

```
[[100, 0.8421537516634251],
 [200, 0.8417442931722796],
 [300, 0.8422561162862114],
 [400, 0.8424608455317842],
 [500, 0.8420513870406388],
 [600, 0.8414371993039206],
 [700, 0.8418466577950661],
 [800, 0.8414371993039206],
 [900, 0.8415395639267069],
 [1000, 0.8414371993039206]]
```

```
score_list_rfc_2=[]
for max_feature in ['auto','sqrt']:
    rfc=
    RandomForestClassifier(random_state=0,max_depth=17,n_estimators=400,max
x_features=max_feature)
    rfc.fit(X_train,y_train)
    y_predict = rfc.predict(X_test)
    accuracy = accuracy_score(y_test,y_predict)
    score_list_rfc_2.append([max_feature,accuracy])
```

score_list_rfc_2

```

[['auto', 0.8424608455317842], ['sqrt', 0.8424608455317842]]

final_rfc_model=
RandomForestClassifier(random_state=0,max_depth=17,n_estimators=400,)
final_rfc_model.fit(X_train,y_train)
y_predict = final_rfc_model.predict(X_test)
accuracy = accuracy_score(y_test,y_predict)

accuracy

0.8424608455317842

columns=['age', 'workclass', 'fnlwgt', 'education', 'education-num',
'marital-status', 'occupation', 'relationship', 'race', 'sex',
'capital-gain', 'capital-loss', 'hours-per-week', 'native-country']
inp=[76, ' Private',124191, ' Masters',14, ' Married-civ-spouse', '
Exec-managerial', ' Husband', ' White', ' Male',0,0,40, ' United-
States']

inp_data={}
i=0

for col in columns:
    inp_data[col]=inp[i]
    i=i+1
inp_data=pd.DataFrame(inp_data)
inp_data=pd.get_dummies(inp_data)
print(type(inp_data))
df_modified_inp=pd.DataFrame()
for col in df_normalized.columns:
    df_modified_inp[col]=[0]
for col in inp_data:
    df_modified_inp[col]=inp_data[col]
df_modified_inp['capital-gain']=df_modified_inp['capital-gain']+5
df_modified_inp['capital-loss']=df_modified_inp['capital-loss']+5

df_modified_inp.drop(columns=['income_ >50K', 'income_
<=50K'],inplace=True)

y_predict = final_rfc_model.predict(df_modified_inp)
print(y_predict)

<class 'pandas.core.frame.DataFrame'>
[1]

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:15:
PerformanceWarning: DataFrame is highly fragmented. This is usually
the result of calling `frame.insert` many times, which has poor
performance. Consider joining all columns at once using
pd.concat(axis=1) instead. To get a de-fragmented frame, use

```



```
`newframe = frame.copy()`  
    from ipykernel import kernelapp as app
```